

WIKIPEDIA & FANDOM ARTICLE RECOMMENDER SYSTEM

A content-based recommender system that suggests similar Wikipedia or Fandom articles based on previously visited pages

Monika Dominiak 160307
Antonina Grzielewska 160286

1. Data Collection: Crawling and Scraping

Article data (2000 articles) was collected directly from Wikipedia and Fandom Wiki using a custom spider written in Scrapy (my_spider.py).

This spider automatically follows internal article links, downloads content, and saves each page's URL, title, and cleaned text to a CSV file (data/wiki.csv or data/fandom.csv).

Key features

Article Validation	Ensures only valid article URLs are scraped
Text Extraction	Extracts <p> tag text, removes punctuation, normalizes whitespace
Minimum Word Count	Filters out short or non-content pages
Random Fallback	When stuck, fetches a random article (Special:Random) to continue exploration
CSV Export	Saves url, title, and text to csv file

Parameters

<code>max_links</code>	Maximum number of pages to collect
<code>start_url</code>	Starting article URL
<code>min_word_count</code>	Minimum number of words to save a page
<code>allow_random</code>	Enables fetching random pages when stuck
<code>user_source_type</code>	Manually sets source to "wiki" or "fandom" for fetching random pages

2. Text Preprocessing

TF-IDF similarity

- Text is preprocessed before vectorization:


Steps applied using NLTK and regex:

1. Lowercasing
2. Stopword removal (`stopwords.words('english')`)
3. Filtering non-alphabetic tokens
4. Lemmatization and stemming comparison

Embedding similarity

- Raw text is generally used without preprocessing.

Database statistics



3. Processing User Input and Selecting Diverse Articles

When a user provides a list of previously visited articles, the system needs to process them in a way that maximizes recommendation quality while keeping computation efficient. This step involves scraping, preprocessing, and selecting the most diverse subset of articles.

- The system first scrapes the user-provided URLs using a simplified scraping function (`run_scraper_simple`).
- For each URL:
 - The system checks whether it is a valid article (`is_article`).
 - It extracts the text and title.
 - Articles below a minimum word count (e.g., 10 words) are discarded.
 - The result is a dataframe of valid user articles, ready for preprocessing.

Similar to the main article database, each user-provided article undergoes:

- Lowercasing
- Stopword removal
- Lemmatization using WordNetLemmatizer

This ensures that user articles are in the same clean, normalized format as the database articles.

Problem:

Users may provide too many articles. Including all articles could:

- Slow down similarity computations
- Over-represent some topics if multiple similar articles are provided

Solution:

Select a subset of the most diverse articles based on their semantic embeddings.

- 1.Embed articles: Convert all user articles into dense sentence embeddings using SentenceTransformer.

- 2.Cluster embeddings: Apply KMeans clustering to group articles into `top_n` clusters (20).

- 3.Select representatives: For each cluster, pick the article closest to the cluster centroid.

- 4.Use these diverse articles as input for the recommendation engine.

4. Feature Representation

In this project, we use two main types of features: TF-IDF vectors (lexical features) and sentence embeddings (semantic features), combined into a hybrid representation.

TF-IDF Representation

- Implements TfidfVectorizer from scikit-learn.
- Captures word-level importance relative to the corpus.
- All articles from a user are combined into one long, normalized, and lemmatized text
- Highlights important words in each article.
- Reduces the weight of common words.
- Works well for matching articles that share similar vocabulary.

Sentence Embeddings

- Uses pretrained SentenceTransformer (all-MiniLM-L6-v2).
- Encodes entire documents into dense vectors representing semantic meaning.
- User input (multiple articles) is combined as a weighted average (later articles are weighted slightly higher for recency).
- Captures semantic similarity beyond exact word matches.
- Handles synonyms, paraphrasing, and conceptual similarity.
- Essential for recommending articles from different topics that are conceptually related.

Hybrid Representation

Both methods are combined to balance textual and conceptual similarity:

Final Score = $0.4 * \text{tfidf_sim} + 0.6 * \text{embedding_sim}$

5. Recommendation algorithm

To make the recommendation process more personalized, the system can take into account the user's reading history — i.e., the articles that the user has already viewed or interacted with. The most recent views receive higher weights to better capture current interests. This is a function parameter, so user can choose which way he wants to do this. We were checking everything taking weighted history into consideration.

This history is used as a context vector that influences which new articles are recommended. The most recently seen articles are the most important.

- We take into consideration only articles that user did not see before
- Preprocess text using the same pipeline.
- Compute TF-IDF and embedding representations.
- Calculate similarity with all articles in the dataset.
- Rank by hybrid score
- Return articles above min_similarity threshold (0.9) or if there are no such articles, return top 10 results.

Example 1

User urls:

<https://en.wikipedia.org/wiki/Monkey>


<https://en.wikipedia.org/wiki/Tiger>

Recommendations:


1	https://en.wikipedia.org/wiki/Monkeys_in_Chinese_culture	0.866419
2	https://en.wikipedia.org/wiki/Platyrrhini	0.822383
3	https://en.wikipedia.org/wiki/Aotidae	0.797528
4	https://en.wikipedia.org/wiki/Howler_monkey	0.79152
5	https://en.wikipedia.org/wiki/Apes	0.730016
6	https://en.wikipedia.org/wiki/Simians_(Chinese_poetry).	0.722337
7	https://en.wikipedia.org/w/index.php?title=Simians_(Chinese_poetry)&oldid=1310610400	0.722337
8	https://en.wikipedia.org/wiki/Pitheciidae	0.718461
9	https://en.wikipedia.org/wiki/Colobinae	0.694462
10	https://en.wikipedia.org/wiki/Patasola_magdalenae	0.680305

For comparison, order_importance = False


No.	title	URL	Similarity
1	Kangchenjunga	https://en.wikipedia.org/wiki/Kangchenjunga	0.793849
2	Ronald Tompkins	https://disney.fandom.com/wiki/Ronald_Tompkins	0.710561
3	Nilgiri Mountains	https://en.wikipedia.org/wiki/Nilgiri_Mountains	0.681216
4	Cricket	https://en.wikipedia.org/wiki/Cricket	0.668438
5	Idiyappam	https://en.wikipedia.org/wiki/String_hoppers	0.662858
6	Rasgulla	https://en.wikipedia.org/wiki/Rasgulla	0.640722
7	Kabaddi	https://en.wikipedia.org/wiki/Kabaddi	0.631143
8	Sulaiman Mountains	https://en.wikipedia.org/wiki/Sulaiman_Mountains	0.619976
9	Kashmir Valley	https://en.wikipedia.org/wiki/Kashmir_Valley	0.614973
10	Geography of West Bengal	https://en.wikipedia.org/wiki/Geography_of_West_Bengal	0.613574




Article 1 breakdown




Article 2 breakdown





Article 3 breakdown




Article 4 breakdown



Article 5 breakdown





We can see that most of recommended articles are near the 'Monkey' article, that means they are related

Example 2 - more diverse urls

User urls:

https://en.wikipedia.org/wiki/Pozna%C5%84_University_of_Technology

https://harrypotter.fandom.com/wiki/Hermione_Granger


https://en.wikipedia.org/wiki/Mount_Everest

<https://en.wikipedia.org/wiki/Doughnut>


<https://en.wikipedia.org/wiki/Volleyball>

Recommendations:


1	https://en.wikipedia.org/wiki/Kangchenjunga	0.683098
2	https://en.wikipedia.org/wiki/Cricket	0.642806
3	https://disney.fandom.com/wiki/Ronald_Tompkins	0.638422
4	https://en.wikipedia.org/wiki/Kabaddi	0.610859
5	https://en.wikipedia.org/wiki/Nilgiri_Mountains	0.59479
6	https://en.wikipedia.org/wiki/String_hoppers	0.593738
7	https://en.wikipedia.org/wiki/Rasgulla	0.574868
8	https://en.wikipedia.org/wiki/Traditional_games_of_South_Asia	0.562681
9	https://en.wikipedia.org/wiki/Kashmir_Valley	0.536844
10	https://en.wikipedia.org/wiki/Sulaiman_Mountains	0.534404




Article 1 breakdown




Article 2 breakdown




Article 3 breakdown



Article 4 breakdown



Article 5 breakdown



Example 3 - very domain-specific

User urls:

https://en.wikipedia.org/wiki/COVID-19_pandemic

<https://en.wikipedia.org/wiki/Vaccine>

<https://en.wikipedia.org/wiki/Virus>

Recommendations:

1	https://en.wikipedia.org/wiki/Virus_classification	1
---	---	---

The result for one article is above our threshold, so it is highly recommended

